

Programming for Data Analytics

Lecture 8: stringr

Dr. Jim Duggan,
School of Engineering & Informatics
National University of Ireland Galway.

https://twitter.com/_jimduggan



Lecture Overview

- stringr basics
- Regular Expressions overview
 - str_view()
 - Anchors
 - Repetition
- Tools
 - str_detect()
 - str_count()
 - str_extract()
 - str_replace()
 - str_split()

Advanced R

*Closures – S3 – S4 – RC Classes –
R Packages – RShiny*

Data Science

*ggplot2 – dplyr – tidyr – stringr – lubridate –
Case Studies*

Base R

*Vectors – Functions – Lists – Matrices –
Data Frames – Apply Functions*



(1) Stringr Basics

- Focus on the stringr package, which is not part of the tidyverse.
- All stringr functions start with **str_**
- Strings can be created with single or double quotes, recommend to use ""
- Multiple strings can be stored in a character vector

```
> s1 <- "This is a string"
> s1
[1] "This is a string"
>
> s2 <- 'This is a string too'
> s2
[1] "This is a string too"
>
> s3 <- c("One", "Two three")
> s3
[1] "One"          "Two three"
```

String Length: `str_length()`

- Returns the number of characters in a string
- Is vectorised

```
> s1
[1] "This is a string"
>
> str_length(s1)
[1] 16
>
> s3 <- c("One", "Two three")
>
> str_length(s3)
[1] 3 9
```

Combining Strings: `str_c()`

- To combine two or more strings, use `str_c()`
- Use the `sep` argument to control how they are separated
- `str_c()` is vectorised, and automatically recycles shorter vectors to the same length as the longest

```
> str_c("x", "y")
[1] "xy"
>
> str_c("x", "y", "z")
[1] "xyz"
>
> str_c("x", "y", "z", sep="-")
[1] "x-y-z"
>
> str_c(c("x", "y", "c"), "ab")
[1] "xab" "yab" "cab"
>
> str_c(c("x", "y", "c"), "ab",
+       collapse = "")
[1] "xabyabcab"
```

Subsetting Strings: `str_sub()`

- As well as the string, this function takes the start and end arguments that give the (inclusive) position of the string.

```
> x <- c("Apple", "Banana", "Pear")
> str_sub(x, 1, 3)
[1] "App" "Ban" "Pea"
>
> str_sub(x, -3, -1)
[1] "ple" "ana" "ear"
```

(2) Matching Patterns with Regular Expressions

- Regexps are very terse language that allows you to describe patterns in strings
- Very powerful features to process strings
- To learn regular expressions:
 - use `str_view()` and `str_view_all()`
 - These take a character vector, and a regular expression and show the match

Basic matches

- The simplest patterns match exact strings

```
x <- c("Apple", "Banana", "Pear")  
str_view(x, "an")
```

Apple
Banana
Pear

- The next step up in complexity is ., which matches any character (except a newline)

```
x <- c("Apple", "Banana", "Pear")  
str_view(x, ".a.")
```

Apple
Banana
Pear

To match ., we need \\.

```
y <- c("abc", "a.c", "bef")  
str_view(y, "a\\.c")
```

abc

a.c

bef

Anchors

- By default, a regular expression will match any part of a string.
- It can be useful to *anchor* the expression so that it matches from the start or end of the string
 - ^ match start
 - \$ match end

```
x <- c("apple pie", "apple", "apple cake")  
str_view(x, "^apple")  
str_view(x, "^apple$")
```

```
apple pie  
apple  
apple cake
```

```
apple pie  
apple  
apple cake
```

Challenge 8.1

- Take a random sample (50, seed=99) from the corpus of common words in `stringr::words` and create regular expressions that find all words that:
 - Start with “w”
 - End with “o”
 - Are exactly three characters long
- The `match` argument can be used to `str_view()` to show the matching/non-matching words.

Other useful matching tools

- `\d` – matches any digit
- `\s` – matches any whitespace
- `[abc]` – matches a, b or c
- `[^abc]` – matches anything except a, b or c.

```
str_view(y, "\\d")      abc  
                        abc1
```

```
str_view(c("grey", "gray"), "gr(ela)y")
```

```
grey  
gray
```

Challenge 8.2

- With the sample sample data set, create regular expressions to find words that:
 - Start with a vowel
 - Start with two successive consonants
 - End with **ing**



Repetition

- Next, we can control how many times a pattern matches
 - ? → 0 or 1
 - + → 1 or more
 - * → 0 or more

```
x <- "AACABCDDD"
```

AACABCDDD

```
str_view(x, "CD?")
```

AACABCD

```
str_view(x, "CD+")
```

```
str_view(x, "CD*")
```

AACABCDDD

Specifying the number of matches

- The number of matches can be specified precisely (default match is “greedy” – will match the longest string possible)
 - {n}: exactly n
 - {n,}: n or more
 - {,m}: at most m
 - {n,m}: between n and m

AACABCDDD

AACABCDDD

AACABCDDD

AACABCDDD

```
x <- "AACABCDDD"
```

```
str_view(x, "A{2}")
```

```
str_view(x, "D{1,}")
```

```
str_view(x, "D{2,3}")
```

```
str_view(x, "D{2,3}?")
```

Challenge 8.3

- Create regular expressions that find all words that:
 - Start with two consonants
 - Have three vowels in a row
- Make use of the sample data set.

(3) Regex Tools in R

- Given the basics just covered, these can now be applied to real problems. The **stringr** package has functions to:
 - Determine which strings match a pattern
 - Find the positions of the matches
 - Extract the content of matches
 - Replace matches with new values
 - Split a string based on a match

Detect Matches: `str_detect()`

- Determines if a character vector matches a pattern, and returns a logical vector of the same length
- Use of `sum()` useful... count number of matches.

```
> x <- c("apple", "banana", "pear")
>
> str_detect(x, "e")
[1] TRUE FALSE TRUE
>
> sum(str_detect(words, "[aeiou]$"))
[1] 271
>
> mean(str_detect(words, "[aeiou]$"))
[1] 0.2765306
```

Using with dplyr and `str_count()`

```
df <- tibble(word=words,i=seq_along(words))
```

```
df %>%
```

```
  mutate(  
    vowels = str_count(word,"[aeiou]"),  
    consonants = str_count(word,"^[^aeiou]")  
  )
```

```
# A tibble: 980 x 4  
      word      i vowels consonants  
  <chr> <int> <int>      <int>  
1     a     1     1          0  
2    able     2     2          2
```

Challenge 8.4

- Use `str_detect()` to find all words that start with a vowel and end in a consonant.
- What word in the corpus has the highest number of vowels?

Exact Matches: `str_extract()`

- To extract the actual text of a match, use `str_extract()`
- Using the Harvard Sentences data set, designed to test VOIP systems

```
> head(sentences)
```

```
[1] "The birch canoe slid on the smooth planks."
[2] "Glue the sheet to the dark blue background."
[3] "It's easy to tell the depth of a well."
[4] "These days a chicken leg is a rare dish."
[5] "Rice is often served in round bowls."
[6] "The juice of lemons makes fine punch."
```

Example

- The task is to locate all sentences that contain a colour.
- `str_subset()` useful as it filters the data set based on a regular expression

```
> colours <- c(
+   "red", "orange", "yellow", "green", "blue", "purple"
+ )
>
> col_match <- str_c(colours, collapse = "|")
>
> col_match
[1] "red|orange|yellow|green|blue|purple"
```

Note: `str_extract()` returns 1st match

```
> has_colour <- str_subset(sentences, col_match)
> matches <- str_extract(has_colour,col_match)
> head(matches)
[1] "blue" "blue" "red"  "red"  "red"  "blue"
>
> more <- sentences[str_count(sentences,col_match) > 1]
> more
[1] "It is hard to erase blue or red ink."
[2] "The green light in the brown box flickered."
[3] "The sky in the west is tinged with orange red."
>
> str_extract(more,col_match)
[1] "blue"    "green"   "orange"
```



```
> str_extract_all(more,col_match)
```

```
[[1]]
```

```
[1] "blue" "red"
```

```
[[2]]
```

```
[1] "green" "red"
```

```
[[3]]
```

```
[1] "orange" "red"
```

```
>
```

```
> str_extract_all(more,col_match,simplify = T)
```

```
      [,1]      [,2]
```

```
[1,] "blue"    "red"
```

```
[2,] "green"    "red"
```

```
[3,] "orange"   "red"
```


Grouped Matches

- Parentheses can be used to clarify precedence
- They can also be used to extract parts of a complex match
- For example, if we need to extract nouns from sentences
- A heuristic:
 - *Any word that come after “a” or “the”*
 - *A word is defined as a sequence of at least one character that isn’t a space*

Example: Two groups

```
noun <- "(a|the) ([^ ]+)"
```

```
has_noun <- sentences %>%  
  str_subset(noun) %>%  
  head(10)
```

```
> has_noun
```

```
[1] "The birch canoe slid on the smooth planks."  
[2] "Glue the sheet to the dark blue background."  
[3] "It's easy to tell the depth of a well."  
[4] "These days a chicken leg is a rare dish."  
[5] "The box was thrown beside the parked truck."
```

Using `str_extract()`

> `has_noun`

```
[1] "The birch canoe slid on the smooth planks."  
[2] "Glue the sheet to the dark blue background."  
[3] "It's easy to tell the depth of a well."  
[4] "These days a chicken leg is a rare dish."  
[5] "The box was thrown beside the parked truck."
```

> `has_noun %>%`

+ `str_extract(noun)`

```
[1] "the smooth" "the sheet"  "the depth"  "a chicken"  
[5] "the parked" "the sun"    "the huge"   "the ball"  
[9] "the woman"  "a helps"
```

Using `str_match()`

- Column 1 is the complete match
- Successive columns are for each group.
- Suggested nouns are in group 3
- *All matches not included*

```
> has_noun %>%  
+   str_match(noun)  
      [,1]      [,2] [,3]  
[1,] "the smooth" "the" "smooth"  
[2,] "the sheet"  "the" "sheet"  
[3,] "the depth"  "the" "depth"  
[4,] "a chicken"  "a"   "chicken"  
[5,] "the parked" "the" "parked"  
[6,] "the sun"     "the" "sun"  
[7,] "the huge"    "the" "huge"  
[8,] "the ball"    "the" "ball"  
[9,] "the woman"   "the" "woman"  
[10,] "a helps"    "a"   "helps"
```

Using with dplyr & extract()

```
tibble(sentence=sentences) %>%  
  tidyr::extract(  
    sentence, c("article", "noun"), "(al the) ([^ ]+)",  
    remove=F  
  )
```

```
# A tibble: 720 x 3
```

	sentence	article	noun
*	<chr>	<chr>	<chr>
1	The birch canoe slid on the smooth planks.	the	smooth
2	Glue the sheet to the dark blue background.	the	sheet
3	It's easy to tell the depth of a well.	the	depth
4	These days a chicken leg is a rare dish.	a	chicken
5	Rice is often served in round bowls.	<NA>	<NA>

Replacing Matches

- `str_replace()` and `str_replace_all()` allow the replacement of matches with new strings.

```
> x <- c("apple", "pear", "banana")
>
> x
[1] "apple"  "pear"   "banana"
>
> str_replace(x, "[aeiou]", "_")
[1] "_pple"  "p_ar"   "b_nana"
>
> str_replace_all(x, "[aeiou]", "_")
[1] "_ppl_"  "p__r"   "b_n_n_"
```

Perform multiple replacements

```
> x <- c("1 house", "2 cars", "3 people")  
>  
> str_replace_all(x, c("1"="one", "2"="two", "3"="three"))  
[1] "one house"      "two cars"       "three people"
```

Flip the order of second and first words (note groupings)

```
> sentences %>%  
+   str_replace("([ ^ ]+) ([ ^ ]+)", "\\2 \\1") %>%  
+   head(5)  
[1] "birch The canoe slid on the smooth planks."  
[2] "the Glue sheet to the dark blue background."  
[3] "easy It's to tell the depth of a well."  
[4] "days These a chicken leg is a rare dish."  
[5] "is Rice often served in round bowls."
```


Challenge 8.5

- Switch the first and last words in the sentences

> s

```
[1] "The birch canoe slid on the smooth planks."  
[2] "Glue the sheet to the dark blue background."  
[3] "It's easy to tell the depth of a well."  
[4] "These days a chicken leg is a rare dish."  
[5] "Rice is often served in round bowls."
```

> ans

```
[1] "planks birch canoe slid on the smooth The."  
[2] "background the sheet to the dark blue Glue."  
[3] "well easy to tell the depth of a It's."  
[4] "dish days a chicken leg is a rare These."  
[5] "bowls is often served in round Rice."
```

Splitting Strings – `str_split()`

- Use `str_split()` to split a string into pieces
- With `simplify = TRUE`, a matrix is returned

```
> x <- "This is a sentence"
```

```
>
```

```
> x
```

```
[1] "This is a sentence"
```

```
>
```

```
> str_split(x, " ")
```

```
[[1]]
```

```
[1] "This"      "is"        "a"         "sentence"
```



Other types of pattern

- When you use a pattern that's a string, it is automatically wrapped into a call to `regex()`
- `regex()` can use other arguments to control the details of the match

```
> x <- c("APPLE", "aPPle", "Apple")
>
> x
[1] "APPLE" "aPPle" "Apple"
>
> x[str_detect(x, regex("apple", ignore_case=T))]
[1] "APPLE" "aPPle" "Apple"
```

comments = TRUE

- Allows you to use comments and white space to make complex regular expressions more understandable

```
> phone <- regex("
+   \\(?      # optional opening parens
+   (\\d{3})  # area code
+   [- ]?    # optional closing parens, dash, or space
+   (\\d{3})  # another three numbers
+   [-]?     # optional space or dash
+   (\\d{4})  # four more numbers
+   ", comments=TRUE)
>
> str_match("514-791-8111", phone)
      [,1]      [,2]  [,3]  [,4]
[1,] "514-791-8111" "514" "791" "8111"
```

Other Uses of Regular Expressions

- `apropos()` searches all objects available from the global environment.

```
> apropos("str_")
```

[1] "str_c"	"str_conv"	"str_count"
[4] "str_detect"	"str_dup"	"str_extract"
[7] "str_extract_all"	"str_interp"	"str_join"
[10] "str_length"	"str_locate"	"str_locate_all"
[13] "str_match"	"str_match_all"	"str_order"
[16] "str_pad"	"str_replace"	"str_replace_all"
[19] "str_replace_na"	"str_sort"	"str_split"
[22] "str_split_fixed"	"str_sub"	"str_sub<-"
[25] "str_subset"	"str_to_lower"	"str_to_title"
[28] "str_to_upper"	"str_trim"	"str_trunc"
[31] "str_view"	"str_view_all"	"str_wrap"

stringi Package

- `stringr` is built on top of `stringi` package
- `stringr` exposes a minimal set of functions
- `stringi` is designed to be comprehensive, and contains 234 functions (as compared to 42 for `stringr`)
- Exercises, find the `stringi` functions that
 - Count the number of words
 - Find duplicated strings
 - Generate random text

<http://tidytextmining.com/index.html>

Text Mining with R

[Welcome to Text Mining with R](#)

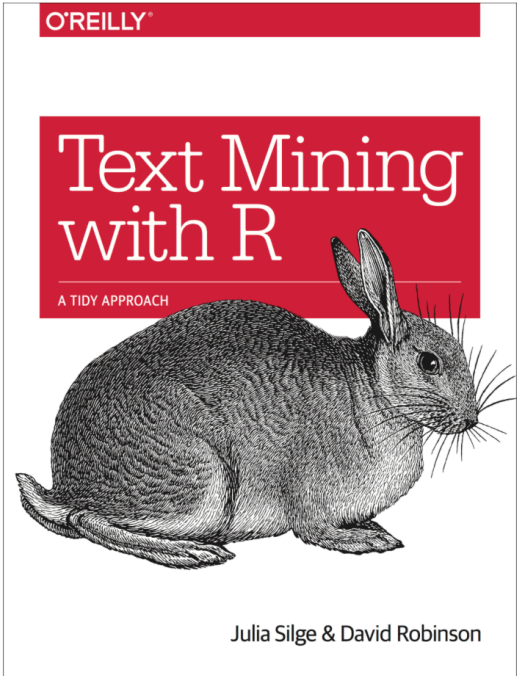
Preface

- 1 The tidy text format
- 2 Sentiment analysis with tidy data
- 3 Analyzing word and document frequency
- 4 Relationships between words: n-grams
- 5 Converting to and from non-tidy format
- 6 Topic modeling
- 7 Case study: comparing Twitter archives
- 8 Case study: mining NASA metadata
- 9 Case study: analyzing usenet text
- 10 References

Published with bookdown

☰ 🔍 A ✎

Welcome to Text Mining with R



This is the [website](#) for *Text Mining with R*! Visit the [GitHub repository](#) for this site.

Summary

- **stringr** package very useful in R
- Regular expressions key to string manipulations
- Rich set of functions to support string processing

Advanced R

*Closures – S3 – S4 – RC Classes –
R Packages – RShiny*

Data Science

*ggplot2 – dplyr – tidyr – stringr – lubridate –
Case Studies*

Base R

*Vectors – Functions – Lists – Matrices –
Data Frames – Apply Functions*

